

Machine learning model solution

- CLUSTERING, CART & RANDOM FOREST



Objectives **greatlearning**

- Business problem: To identify amongst the existing customers of a bank who are more likely to take up personal loan , for effective cross-selling
- Modeling objectives:
 - Clustering : Create meaningful clusters from the population, to understand the nature of the customers
 - Classification Tree model: Find rules to classify the customers for cross-selling, techniques to be used
 - CART
 - Random forest


Exploratory data analysis

- Dataset provided has the following characteristics:
 - No. of records: 5000
 - No. of predictor variables : 13
 - 1 customer identifier column, which do not hold any statistical significance
 - 6 continuous numeric variables
 - 4 ordinal variables with 2 levels each
 - 2 ordinal variables with multiple levels
 - Target variable : with 2 distinct classes , 0 for non-responders and 1 for no response
- 480 of 5000 customers responded positively, i.e. only 9.6% success rate


Data analysis : univariate – ordinal variables



Data analysis : univariate – ordinal variables




Data analysis : univariate – ordinal variables




- Majority of the customers currently have only single relationship with the bank
- Securities and deposit accounts are not very popular with the customers
- Credit Card is the most popular product amongst all
- Customers are open to internet banking facilities

Data analysis : univariate – ordinal variables



- Distribution of customers by their educational qualification is mostly uniform

Data analysis : univariate – ordinal variables




- Distribution of customers by number of family members they have
- Since only 0.36% percent of records have NA will exclude these records from further analysis

Data analysis : multivariate – continuous variables




- There is high degree of correlation between customer age and years of professional experience, which is as per expectation
- Will exclude experience years from next steps of analysis
- People with high income tend to spend more on credit card per month, which is also as per expectation, but correlation is just about 0.4 so will keep both these variables

Data analysis : univariate – continuous variables



- Distribution of Age is normal
- Income and Credit card average spending per month have outliers
- These variables need to be scaled

Data analysis : univariate – continuous variables



- Check on records where home mortgage value is greater than 0 , shows that there are quite a number of outliers

Data analysis : Treatment of zip code

- There are unique 467 zipcodes in the dataset
- Since Zipcode as a number do not have any significance will treat it as a factor
- To reduce the cardinality of Zipcode will aggregate it to State level
- Since majority of the zipcodes are of 5 digits assuming them to be from US , hence using “zipcode” package to map them to States
- 33 records are found to have invalid Zipcode, ignoring them , found only one record belonging to state AE and rest 4966 to CA

	0	1
AE	1	0
CA	4489	477

- So we can safely assume that State is a good proxy for Zipcode as it is more related to target variable hence dropping zipcode column


Data transformations done

- Row filters :
 - Records with family member count not specified (18 of 5000)
- Column Filter :
 - ID
 - Zipcode
 - Professional experience years
- Columns scaled for clustering:
 - Age
 - Income
 - CC Average spent
 - House Mortgage amount
- Columns created:
 - Have multiple relationships with bank (1 : Yes , 0:No)
 - Have House Mortgage (1:Yes , 0 :No)

CLUSTERING

Algorithm choice rationale

- Since the number of records are running into thousands we should be avoiding hierarchical clustering methods
- We have both continuous and factor data with more than one levels available within independent variables, hence Euclidean distance function cannot be used
- Gower distance function will be used, which takes care of both the data types
- K-means algorithms cannot be used since it assumed Euclidean distance , hence PAM (partitioning around medoids) will be used which is a variant of k-means , but instead of always defaulting to Euclidean distance it can refer to custom distance function
- We need to provide the number of clusters to start with as input to this algorithm



profiling clusters – Cluster 1

```
[[1]]
age_in_years    income_in_K_month    CCAvg    Mortgage    family_size_fact    Education_fact    have_securities_acct_fact
Min.   :23.00    Min.   : 8.00    Min.   :0.000    Min.   : 0.00    1:107    1:165    0:1164
1st Qu.:33.00    1st Qu.: 33.00    1st Qu.:0.600    1st Qu.: 0.00    2:131    2:909    1: 137
Median :40.00    Median : 54.00    Median :1.300    Median : 0.00    3:293    3:227
Mean   :42.44    Mean   : 61.68    Mean   :1.551    Mean   : 27.25    4:770
3rd Qu.:52.00    3rd Qu.: 80.00    3rd Qu.:2.100    3rd Qu.: 0.00
Max.   :67.00    Max.   :194.00    Max.   :8.300    Max.   :612.00
have_deposit_account_fact    have_online_access_fact    have_cc_fact    have_multiple_relns_fact    have_mortgage_fact
0:1213    0: 284    0:941    0:1213    0:1093
1: 88    1:1017    1:360    1: 88    1: 208
```

- Characterized by
 - Majority are Graduates
 - Majority have online access
 - Majority have 4 family members
 - Majority do not have mortgage account

profiling clusters – Cluster 2

```
[[2]]
age_in_years    income_in_K_month    CCAvg    Mortgage    family_size_fact    Education_fact    have_securities_acct_fact
Min.   :23.00    Min.   : 8.00    Min.   : 0.000    Min.   : 0.00    1:170                1:272                0:1110
1st Qu.:38.00    1st Qu.: 34.00    1st Qu.: 0.600    1st Qu.: 0.00    2:602                2:212                1: 134
Median :49.00    Median : 58.00    Median : 1.500    Median : 0.00    3:263                3:760
Mean   :47.36    Mean   : 66.19    Mean   : 1.726    Mean   : 25.26    4:209
3rd Qu.:57.00    3rd Qu.: 85.00    3rd Qu.: 2.200    3rd Qu.: 0.00
Max.   :67.00    Max.   :203.00    Max.   :10.000    Max.   :581.00
have_deposit_account_fact    have_online_access_fact    have_cc_fact    have_multiple_relns_fact    have_mortgage_fact
0:1224                        0:1059                        0:869            0:1195                        0:1055
1: 20                          1: 185                        1:375            1: 49                          1: 189
```

- Characterized by
 - Majority have advanced or professional degree
 - Majority do not have internet banking access
 - Majority have 2 family members
 - Majority do not have any house mortgage

profiling clusters – Cluster 3


```
[[3]]
age_in_years    income_in_K_month    CCAvg    Mortgage    family_size_fact    Education_fact    have_securities_acct_fact
Min.   :23.00    Min.   : 8.00    Min.   : 0.000    Min.   :0    1:757    1:1008    0:1166
1st Qu.:37.00    1st Qu.: 49.75    1st Qu.: 1.000    1st Qu.:0    2:269    2: 78    1: 134
Median :45.00    Median : 84.00    Median : 2.000    Median :0    3:205    3: 214
Mean   :45.71    Mean   : 90.79    Mean   : 2.471    Mean   :0    4: 69
3rd Qu.:55.00    3rd Qu.:128.00    3rd Qu.: 3.300    3rd Qu.:0
Max.   :67.00    Max.   :224.00    Max.   :10.000    Max.   :0
have_deposit_account_fact    have_online_access_fact    have_cc_fact    have_multiple_relns_fact    have_mortgage_fact
0:1197    0: 284    0:903    0:1193    0:1300
1: 103    1:1016    1:397    1: 107    1: 0
```

- Characterized by
 - Majority are undergraduates
 - Majority have online access
 - Majority have credit card accounts
 - Majority have family size of 1
 - None of them have any mortgage account

profiling clusters – Cluster 4

```
[[4]]
age_in_years    income_in_K_month    CCAvg    Mortgage    family_size_fact    Education_fact    have_securities_acct_fact
Min.   :23.00    Min.   : 8.00    Min.   :0.000    Min.   : 75    1:430    1:643    0:1023
1st Qu.:36.00    1st Qu.: 38.00    1st Qu.:0.700    1st Qu.:113    2:290    2:200    1: 114
Median :46.00    Median : 64.00    Median :1.500    Median :158    3:248    3:294
Mean   :45.97    Mean   : 76.28    Mean   :2.012    Mean   :189    4:169
3rd Qu.:56.00    3rd Qu.:104.00    3rd Qu.:2.700    3rd Qu.:234
Max.   :66.00    Max.   :205.00    Max.   :9.000    Max.   :635
have_deposit_account_fact    have_online_access_fact    have_cc_fact    have_multiple_relns_fact    have_mortgage_fact
0:1048    0:386    0:804    0:1049    0: 0
1: 89    1:751    1:333    1: 88    1:1137
```

- Characterized by
 - Majority are undergraduates
 - All of them have mortgage account
 - Majority have internet banking access



CLASSIFICATION

Test and train split

```
## 70% of the sample size
sample_size <- floor(0.7 * nrow(cust_base_orig_fact))

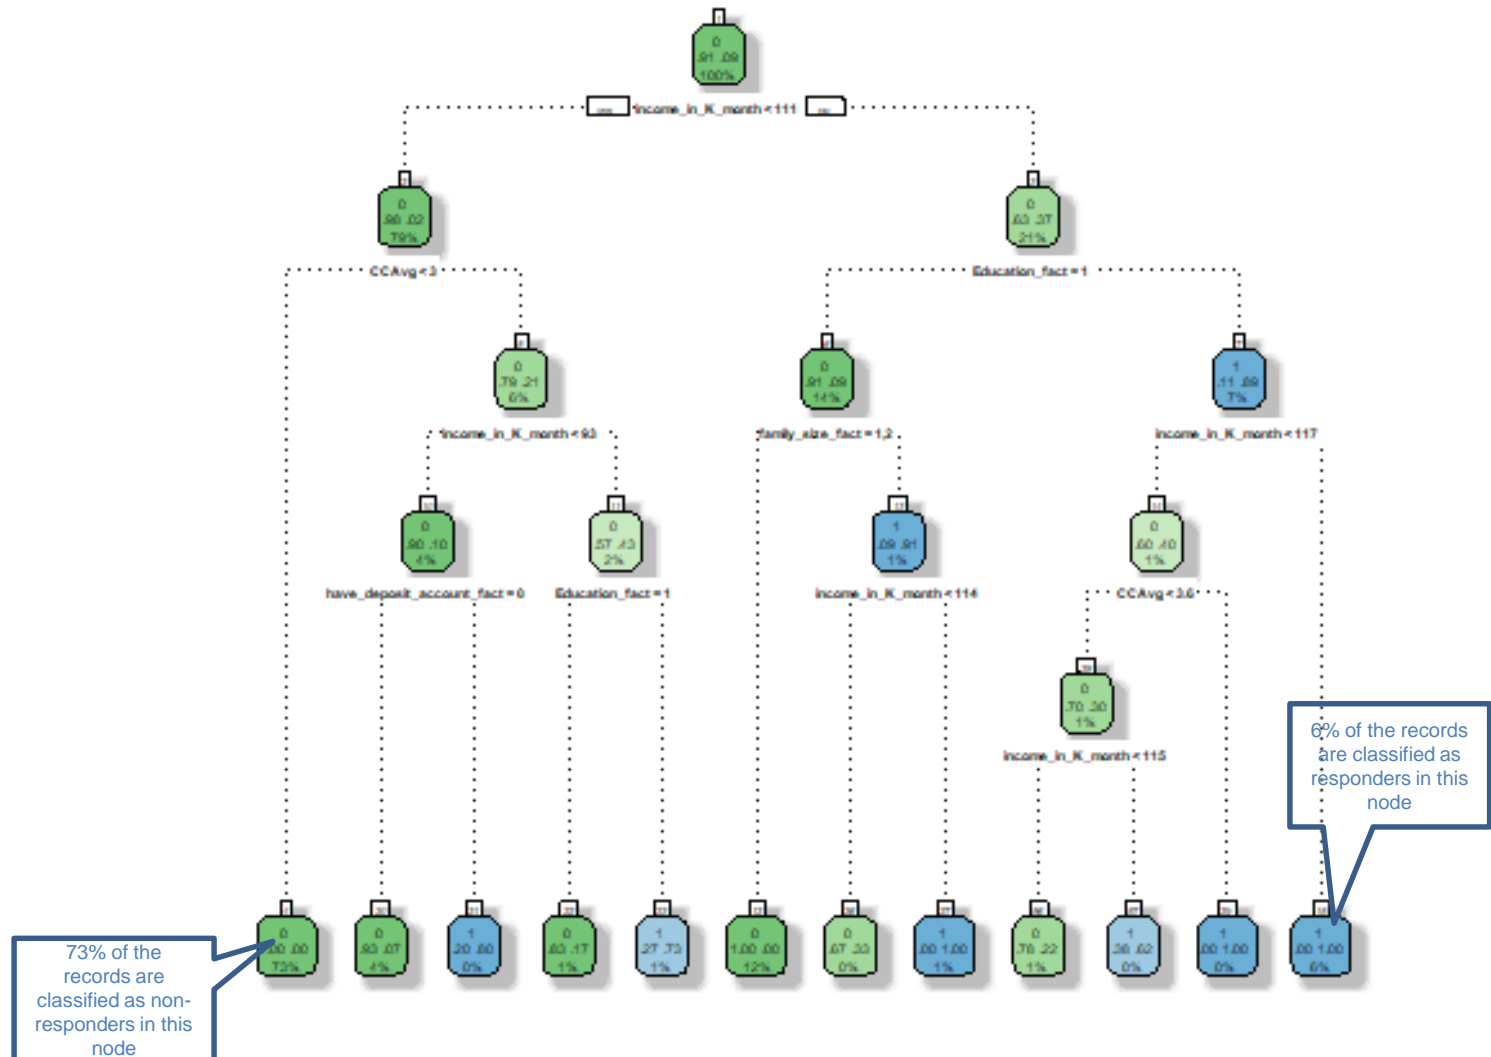
## set the seed to make the partition reproducible
set.seed(4)
train_ind <- sample(seq_len(nrow(cust_base_orig_fact)), size = sample_size)

train_dataset <- cust_base_orig_fact[train_ind, ]
test_dataset <- cust_base_orig_fact[-train_ind, ]

##Check for uniformity of data in test and train
> sum(as.integer(as.character(train_dataset$did_accept_personal_loan_offer_fact))) / nrow(train_dataset)
[1] 0.08976197
> sum(as.integer(as.character(test_dataset$did_accept_personal_loan_offer_fact))) / nrow(test_dataset)
[1] 0.1103679
> sum(as.integer(as.character(cust_base_orig_fact$did_accept_personal_loan_offer_fact))) / nrow(cust_base_orig_fact)
[1] 0.0959454
```

- Data is split into 70:30 ratio for train to test
- Response rate in split and overall data are close

CART



Variables actually used in tree construction:

[1] CCAvg Education_fact family_size_fact have_deposit_account_fact
[5] income_in_K_month

Root node error: 313/3487 = 0.089762

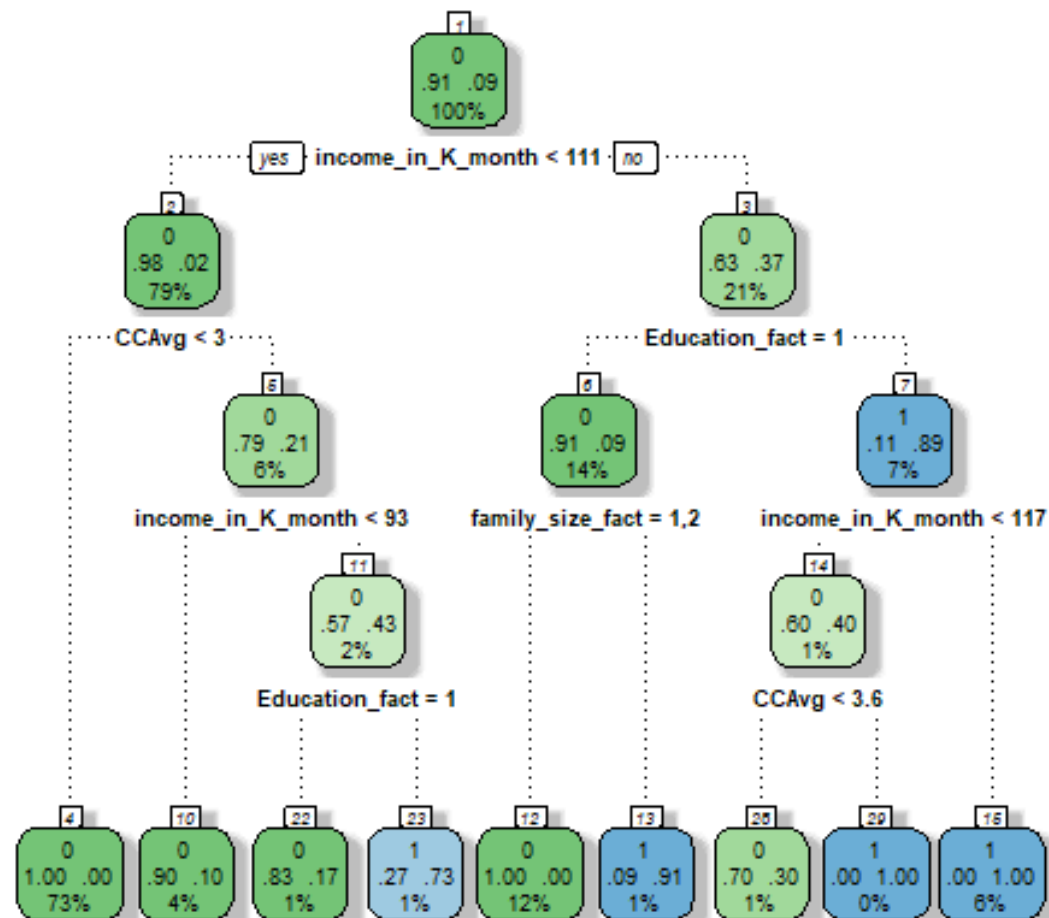
n= 3487

	CP	nsplit	rel error	xerror	xstd
1	0.3210863	0	1.00000	1.00000	0.053927
2	0.1214058	2	0.35783	0.38339	0.034391
3	0.0287540	3	0.23642	0.23962	0.027369
4	0.0223642	4	0.20767	0.23003	0.026828
5	0.0149095	5	0.18530	0.21406	0.025899
6	0.0095847	8	0.14058	0.16613	0.022866
7	0.0063898	9	0.13099	0.17891	0.023716
8	0.0000000	11	0.11821	0.20128	0.025129

Pruned tree

ptree<- prune(tree_iter1, cp= 0.0095847 ,"CP")

Pruned Classification Tree



Model performance measures

- Deciling and rank-ordering the deciles give us the following results
 - More than 90% of the records are accounted for in 10th and 9th deciles
 - KS statistics is 0.93 (benchmark: >0.45) in the first decile itself which indicates overfit
 - Lift calculated is also 9 in the first decile
- Checking on other performance measures:
 - AUC : 0.995953 (benchmark: >0.80)
 - Gini : 0.9028706 (benchmark : >0.60)

	deciles	cnt	cnt_resp	cnt_non_resp	lift
1	10	368	299	69	9.05
2	9	2687	14	2673	1.14
3	2	432	0	432	1.00

Validate against holdout test data

- Running the tree for test dataset created from the initial split gives the following results
 - More than 90% of the records are accounted for in 10th and 9th deciles
 - KS statistics is 0.92 (for train: 0.93) in the first decile itself
 - Lift calculated is also 8.52 in the first decile
- Checking on other performance measures:
 - AUC : 0.9843769 (for train: 0.995953)
 - Gini : 0.8873045 (for train : 0.9028706)
- We see above that the performance measures in train and test are comparable , with not much deviation, hence we can conclude that this is a **good model**.

	deciles	cnt	cnt_resp	cnt_non_resp	lift
1	10	152	143	9	8.52
2	9	1176	22	1154	1.13
3	2	167	0	167	1.00

RANDOM FOREST

Initial model run

Started with these parameters given the guidelines

```
tuned_rf <- tuneRF(x = train_dataset[, -13],
  y=train_dataset$did_accept_personal_loan_offer_fact,
  mtryStart = 3,
  ntreeTry=200,
  stepFactor = 1.5,
  improve = 0.001,
  trace=T,
  plot = T,
  doBest = TRUE,
  nodesize = 15,
  importance=T
)
```

Suggests the out of bag error rate is lowest for mtry=12 so will run the model with that

```
mtry = 3  OOB error = 1.49%
Searching left ...
mtry = 2      OOB error = 1.84%
-0.2307692 0.001
Searching right ...
mtry = 4      OOB error = 1.35%
0.09615385 0.001
mtry = 6      OOB error = 1.26%
0.06382979 0.001
mtry = 9      OOB error = 1.2%
0.04545455 0.001
mtry = 12     OOB error = 1.12%
0.07142857 0.001
```


final model run

Through multiple iterations we see that these parameters give consistent output for multiple runs even on slight variations of the parameters

```
tuned_rf <- tuneRF(x = train_dataset[, -13],
  y=train_dataset$did_accept_personal_loan_offer_fact,
  mtryStart = 8,
  ntreeTry=500,
  stepFactor = 1.5,
  improve = 0.001,
  trace=T,
  plot = T,
  doBest = TRUE,
  nodesize = 200,
  importance=T)
```

Suggests the out of bag error rate is lowest for mtry=8 so will stick to this

```
mtry = 8  OOB error = 1.81%
Searching left ...
mtry = 6      OOB error = 1.86%
-0.03174603 0.001
Searching right ...
mtry = 12     OOB error = 2.15%
-0.1904762 0.001
```

Checking for the importance of variables

	0	1	MeanDecreaseAccuracy	MeanDecreaseGini
Education_fact	45.45	40.05	46.17	145.67
income_in_K_month	33.95	34.27	34.27	123.52
family_size_fact	17.53	15.18	17.61	37.73
CCAvg	12.55	12.42	12.59	31.87
have_deposit_account_fact	5.54	1.39	4.51	10.07
Mortgage	2.13	0.30	2.10	0.49
have_multiple_relns_fact	2.92	-1.73	1.61	1.44
have_CC_fact	1.39	1.42	1.41	0.08
have_securities_acct_fact	-1.00	1.00	1.00	0.01
age_in_years	-1.00	1.00	0.00	0.02
have_online_access_fact	0.00	0.00	0.00	0.00
have_mortgage_fact	0.00	0.00	0.00	0.00

- We see that the 2 new variables added as part of feature engineering process do not have any significance
- Top 4 most important variables helping in reducing impurity in each node is same as in CART model

Model performance measures

- Deciling and rank-ordering the deciles give us the following results
 - More than 100% of the records are accounted for in top 3 deciles
 - KS statistics is 0.94 (benchmark: >0.45) in the first decile itself which indicates overfit
 - Lift calculated is also 9 in the first decile
- Checking on other performance measures:
 - AUC : 0.995788 (benchmark: >0.80)
 - Gini : 0.9005289 (benchmark : >0.60)

	deciles	cnt	cnt_resp	cnt_non_resp	lift
1	10	364	296	68	9.06
2	9	355	16	339	4.83
3	8	2768	1	2767	1.00

Validate against holdout test data

- Running the tree for test dataset created from the initial split gives the following results
 - All 100% of the records are accounted for in top 3 deciles
 - KS statistics is 0.92 (for train: 0.93) in the first decile itself
 - Lift calculated is also 8.42 in the first decile
- Checking on other performance measures:
 - AUC : 0.9809547 (for train: 0.995788)
 - Gini : 0.8822665 (for train : 0.9005289)
- We see above that the performance measures in train and test are comparable , with not much deviation, hence we can conclude that this is a **good model**.

	deciles	cnt	cnt_resp	cnt_non_resp	lift
1	10	150	139	11	8.40
2	9	164	22	142	4.65
3	8	1181	4	1177	1.00

conclusion

- CART and Random Forest have given the similar kind of variable importance output, so following variables play major role in identifying customers who are more likely to respond:
 - Count of Family Members
 - Education Level
 - Income per month
- For both the model performance measurements give us high accuracy indicating overfit, but the model when validated against holdout sample gives comparable measurements
- Imbalanced nature of the data might lead to these observations, though decision trees are theoretically more immune to imbalanced data
- We need to take this model with a pinch of salt, and more data will be good
- Following approaches can be taken to make the model more reliable:
 - Get more data from business
 - Run the models after synthetically correcting for imbalance

THANK YOU